Cencori App Summary

Source: README.md and ARCHITECTURE.md (repo evidence).

What It Is

Cencori is a unified infrastructure platform for AI production that provides security, observability, and scale across AI workloads. It aims to replace one-off scaffolding with a single platform spanning gateway, compute, workflows, storage, and integrations.

Who It's For

Teams building production-grade AI applications or AI businesses that need security, routing, orchestration, and operational visibility without building those layers in-house.

What It Does (Key Features)

- Unified AI gateway with multi-provider routing (OpenAI, Anthropic, Gemini, etc.).
- Built-in security: PII detection, prompt injection protection, and content filtering.
- Observability: audit logs, latency, and cost attribution for requests.
- Streaming responses with token counting and cost tracking.
- Compute for serverless agent execution and edge inference (planned in architecture).
- Workflow orchestration with circuit breakers and state management (planned).

How It Works (Architecture Overview)

- Data flow: Your app -> Cencori platform primitives (Gateway, Compute, Workflow, Storage, Integration) -> AI providers.
- Platform stack: Next.js 15 (App Router) app, TypeScript, Supabase for auth and database, Vercel deployment.
- Entry points (documented): /api/ai/chat live; /api/compute/run, /api/workflow/execute, /api/storage/vectors, /api/integrations/{connector} listed as planned.

How To Run (Minimal Getting Started)

- Prereqs: Node.js 18+, a Supabase project, and AI provider API keys (OpenAI/Anthropic/Google).
- Install: npm install
- Configure: create .env.local with Supabase keys, provider keys, and ENCRYPTION_SECRET.
- Migrations: apply database/migrations/ to Supabase (manual step).
- Start: npm run dev, then open http://localhost:3000.